

Basic RL.6

Judy Tutorial

Softmax (Multiclass Logistic) Regression

&

Neural Networks

MNIST dataset

[LeCun 1998]

Multiple Labels

$y \in \{0,1,2, \dots, 9\}$

Hypothesis function

$$h_k(\theta) = g(\theta_k^T X)$$

inputs
parameters

Activation function

$$g(z_k) = \frac{e^{z_k}}{\sum_{l=1}^K e^{z_l}}$$

softmax

Objective function

$$J(\theta) = -\frac{1}{N} \sum_{i=1}^N \sum_{l=1}^K y_l^{(i)} \log h_l(\theta)^{(i)}$$

Cross entropy loss

Hypothesis function

$$h(\boldsymbol{\theta}) = \begin{bmatrix} g(\boldsymbol{\theta}_{k=1}^T X) \\ g(\boldsymbol{\theta}_{k=2}^T X) \\ g(\boldsymbol{\theta}_{k=3}^T X) \\ \vdots \\ \vdots \\ g(\boldsymbol{\theta}_{k=K}^T X) \end{bmatrix} = \begin{bmatrix} \frac{\exp \boldsymbol{\theta}_{k=1}^T X}{\sum_{l=1}^K \exp \boldsymbol{\theta}_l^T X} \\ \frac{\exp \boldsymbol{\theta}_{k=2}^T X}{\sum_{l=1}^K \exp \boldsymbol{\theta}_l^T X} \\ \frac{\exp \boldsymbol{\theta}_{k=3}^T X}{\sum_{l=1}^K \exp \boldsymbol{\theta}_l^T X} \\ \vdots \\ \vdots \\ \frac{\exp \boldsymbol{\theta}_{k=K}^T X}{\sum_{l=1}^K \exp \boldsymbol{\theta}_l^T X} \end{bmatrix}$$

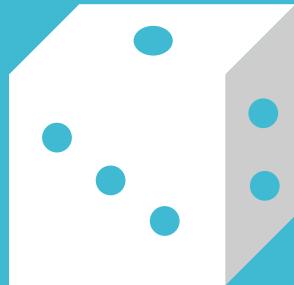
Objective function

$$J(\boldsymbol{\theta}) = -\frac{1}{N} \sum_{i=1}^N \sum_{l=1}^K y_l^{(i)} \log h_l(\boldsymbol{\theta})^{(i)}$$



Negative Log-Likelihood of
Multinomial distribution

Multinomial Distribution



Wiki:

- the Multinomial distribution gives the probability of any particular combination of numbers of successes for the various categories

Example:

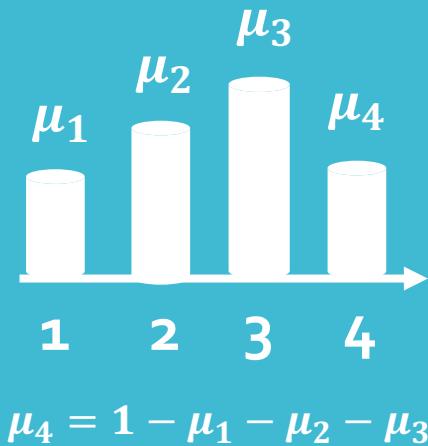
- Repeat rolling a dice for many times
- Suppose we roll a dice 5 times, what's the probability of the appearances of two 3s and three 4s?

$$\cdot p(34344|\mu) = \frac{5!}{2!3!} \left(\frac{1}{6}\right)^2 \left(\frac{1}{6}\right)^3$$



Parameter vector for dice probability
where $\mu_k = \frac{1}{6}$, $K = 6$

Multinomial Distribution



- x_k - the number of repetition of a specific outcome when conducting some events
- μ_k - the probability of each event

$$Mult(x_1, x_2, \dots, x_K | \boldsymbol{\mu}, N) = C_N^{x_1 x_2 \dots x_K} \prod_{k=1}^K \mu_k^{x_k}$$

#events
|
distribution parameter vector

- $0 \leq \mu_k \leq 1, \sum_{k=1}^K \mu_k = 1$
- $\sum_{k=1}^K x_k = N$
- $C_N^{x_1 x_2 \dots x_K} = \frac{N!}{x_1! x_2! \dots x_K!}$

$$Mult(x_1, x_2, \dots, x_K | \boldsymbol{\mu}) \propto \prod_{k=1}^K \mu_k^{x_k}$$

Multinomial Likelihood

#data

$$\mathcal{L}(\boldsymbol{\theta}) = \prod_{i=1}^N \prod_{l=1}^K$$

$$h_l(\boldsymbol{\theta})^{(i)} y_l^{(i)}$$

$$g(\boldsymbol{\theta}_k^T X)$$

$$y_l^{(i)} = \mathbf{1}_{y=k} = \begin{cases} 1, & \text{if } y = k \\ 0, & \text{if } y \neq k \end{cases}$$

One hot
classification
label

$$0 \leq h_k(\boldsymbol{\theta}) \leq 1$$

$$\sum_l^K h_l(\boldsymbol{\theta}) = 1$$

MLE for Softmax Regression

$$\max_{\theta} \log \mathcal{L}(\theta) = \sum_{i=1}^N \sum_{l=1}^K y_l^{(i)} \log h_l(\theta)^{(i)}$$



$$\min_{\theta} J(\theta) = -\frac{1}{N} \sum_{i=1}^N \sum_{l=1}^K y_l^{(i)} \log h_l(\theta)^{(i)}$$

Cross entropy loss

Objective function

$$\min_{\theta} J(\theta) = -\frac{1}{N} \sum_{i=1}^N \sum_{l=1}^K y_l^{(i)} \log h_l(\theta)^{(i)}$$

Batch Gradient

$$\begin{aligned}\nabla_{\theta_j} J(\theta_j) &= \nabla_{\theta_j} - \frac{1}{N} \sum_{i=1}^N \sum_{l=1}^K y_l^{(i)} \log h_l(\theta)^{(i)} \\ &= -\frac{1}{N} \sum_{i=1}^N \sum_{l=1}^K \frac{y_l^{(i)}}{h_l(\theta_j)^{(i)}} \nabla_{\theta_j} h(\theta_j)^{(i)} \\ &= \frac{1}{N} \sum_{i=1}^N \sum_{l=1}^K [h(\theta_j)^{(i)} - y_l^{(i)}] x^{(i)}\end{aligned}$$

Element-wise

Matrix-wise

$$= \frac{1}{N} X[h(\theta) - Y]^T$$

same as linear/logistic regression
related to Exponential Family
and Generalized Linear Model

Gradient Descent $\theta_{t+1} \leftarrow \theta_t - \alpha \nabla_{\theta} J(\theta)$

Batch

Require a batch/entire dataset

$$\nabla_{\theta_j} J(\theta_j) = \frac{1}{N} \sum_{i=1}^N [h(\theta_j)^{(i)} - y^{(i)}] x^{(i)}$$

Sample random mini batch from the entire dataset

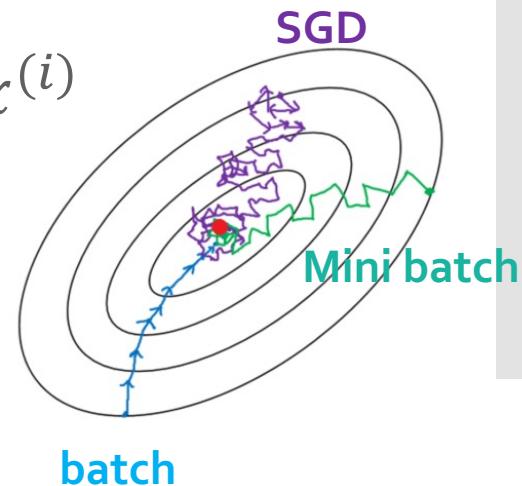
$$\nabla_{\theta_j} J(\theta_j) = \frac{1}{N_{mini}} \sum_{i=1}^{N_{mini}} [h(\theta_j)^{(i)} - y^{(i)}] x^{(i)}$$

&

Stochastic
Gradient
Descent

Require only one data point

$$\nabla_{\theta_j} J(\theta_j) = [h(\theta_j)^{(i)} - y^{(i)}] x^{(i)}$$



Feedforward Neural Networks

(MultiLayer Perceptron)

$$f(\textcolor{blue}{X}) = f^{[M]} \dots (f^{[3]}(f^{[2]}(f^{[1]}(\textcolor{blue}{X}))))$$

First layer
Second layer
Third layer
Mth layer

$$h(\theta) = g(\theta^T X)$$



$$h(W, b) = g(XW + b)$$

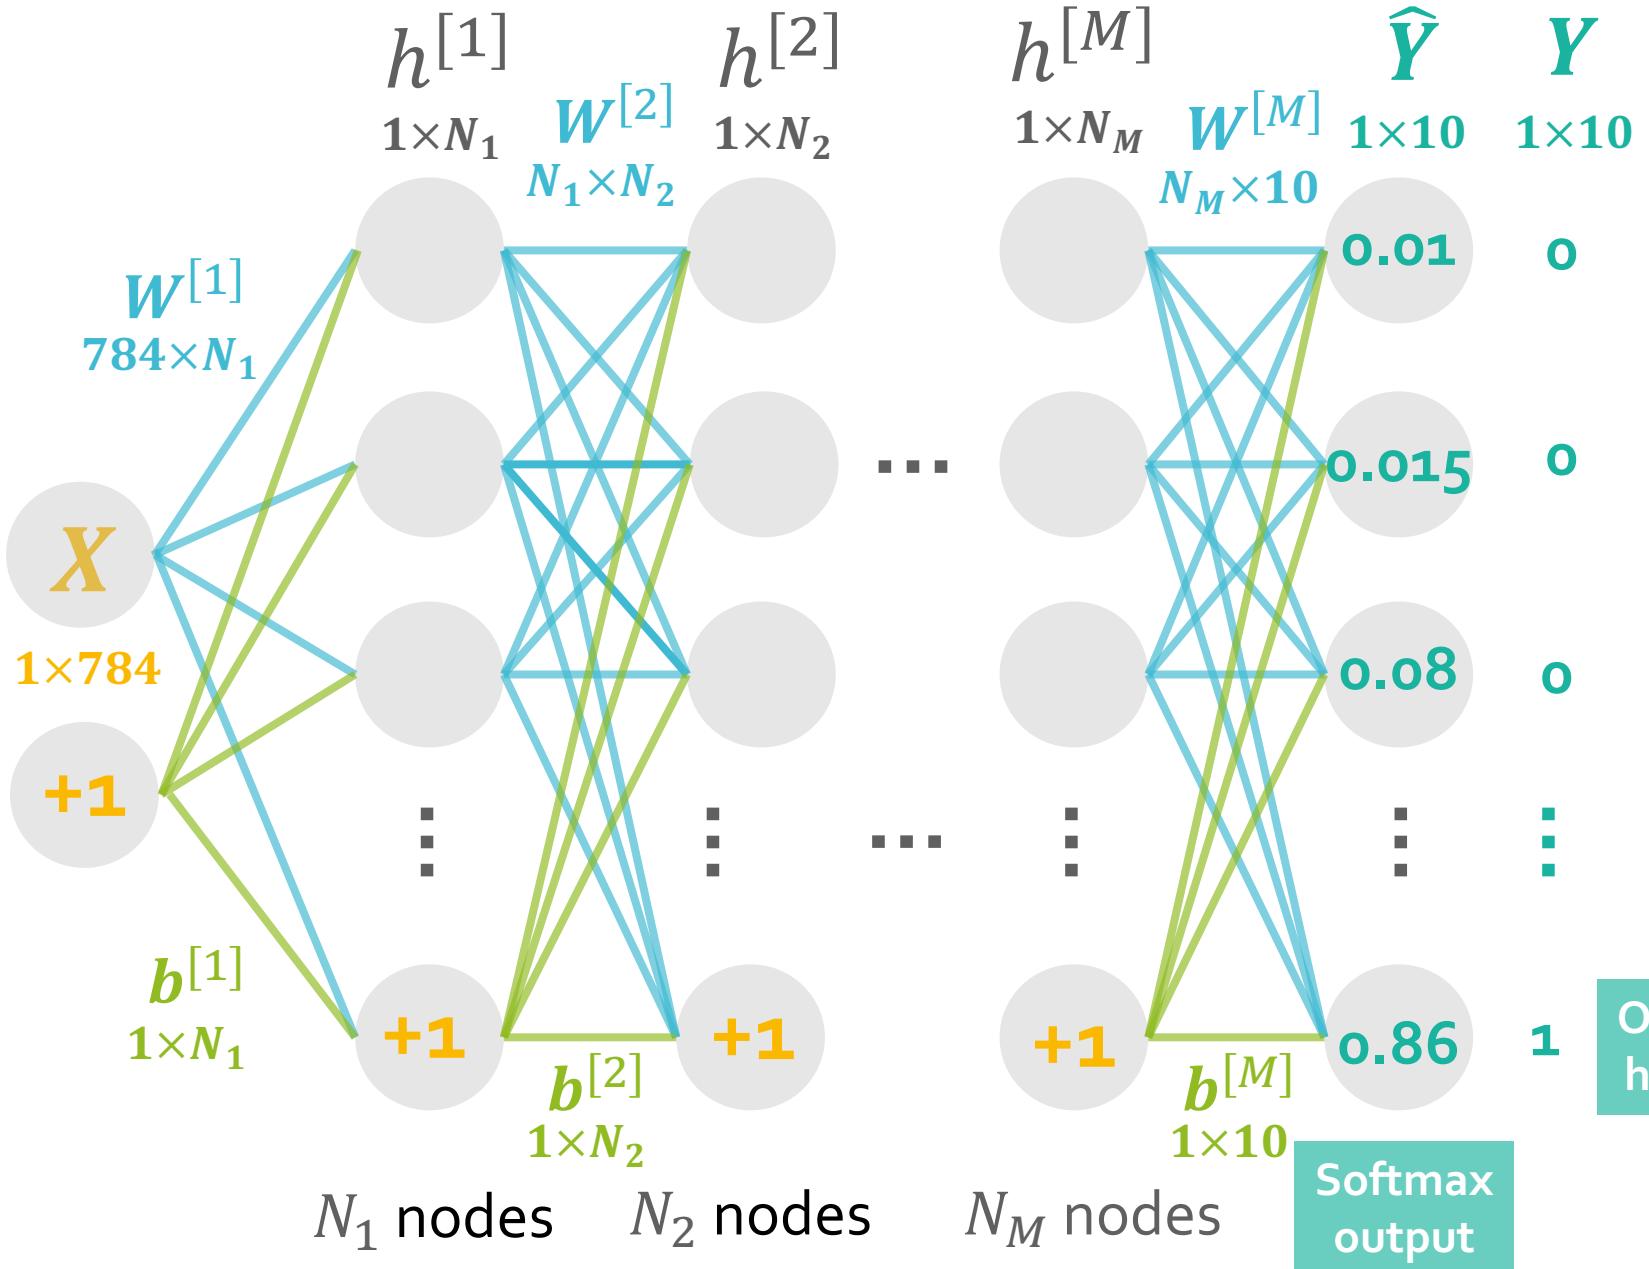
inputs ↗ ↘ parameters

$$h^{[1]}(W^{[1]}, b^{[1]}) = g^{[1]}(XW^{[1]} + b^{[1]})$$

$$h^{[2]}(W^{[2]}, b^{[2]}) = g^{[2]}(h^{[1]}W^{[2]} + b^{[2]})$$

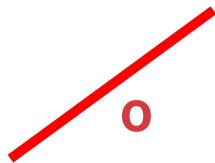
$$h^{[3]}(W^{[3]}, b^{[3]}) = g^{[3]}(h^{[2]}W^{[3]} + b^{[3]})$$

⋮



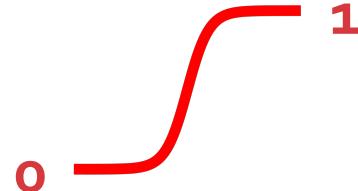
Activation functions

$$g(z) = z$$



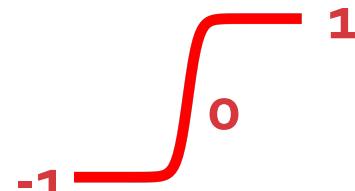
linear

$$g(z) = \frac{1}{1 + e^{-z}}$$



Sigmoid

$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$



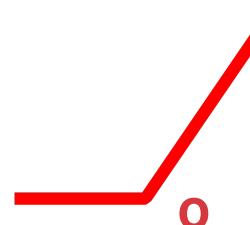
tanh

$$g(z_k) = \frac{e^{z_k}}{\sum_{l=1}^K e^{z_l}}$$

A generalized sigmoid

softmax

$$g(z) = \max(z, 0)$$



ReLU

$$g(z) = \max(z, 0.1z)$$



Leaky ReLU

Objective functions

$$\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

MSE

$$-[y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$$

Binary Cross entropy

$$\frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$

MAE

$$-\frac{1}{N} \sum_{i=1}^N y_i \log \hat{y}_i$$

Multiclass Cross entropy

Optimizers

Batch Gradient

MiniBatch Gradient

SGD

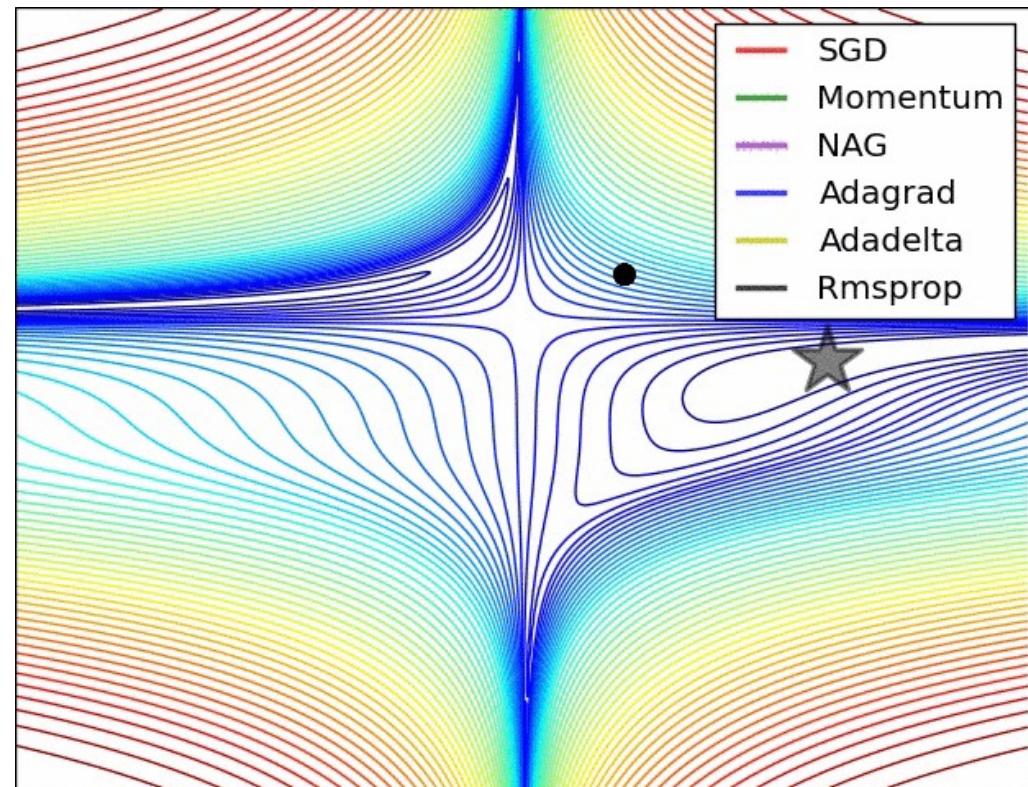
SGD+momentum

Adam

Adagrad

Rmsprop

...



<https://www.analyticsvidhya.com/blog/2021/10/a-comprehensive-guide-on-deep-learning-optimizers/>

Backpropagation – chain rule

$$h^{[1]} = g^{[1]}(XW^{[1]} + b^{[1]})$$

ReLU

$$h^{[2]} = g^{[2]}(h^{[1]}W^{[2]} + b^{[2]})$$

ReLU

$$\hat{Y} = g^{[3]}(h^{[2]}W^{[3]} + b^{[3]})$$

softmax

$$J = -\sum_{i=1}^N y_i \log \hat{y}_i \rightarrow -\log \hat{y}_i$$

Multiclass Cross entropy

$h^{[2]}$

$$\frac{\partial J}{\partial W^{[3]}} = \frac{\partial J}{\partial \hat{Y}} \frac{\partial \hat{Y}}{\partial g^{[3]}} \frac{\partial g^{[3]}}{\partial W^{[3]}}$$

$h^{[1]}$

$$\frac{\partial J}{\partial W^{[2]}} = \frac{\partial J}{\partial \hat{Y}} \frac{\partial \hat{Y}}{\partial g^{[3]}} \frac{\partial g^{[3]}}{\partial h^{[2]}} \frac{\partial h^{[2]}}{\partial g^{[2]}} \frac{\partial g^{[2]}}{\partial W^{[2]}}$$

X

$$\frac{\partial J}{\partial W^{[1]}} = \frac{\partial J}{\partial \hat{Y}} \frac{\partial \hat{Y}}{\partial g^{[3]}} \frac{\partial g^{[3]}}{\partial h^{[2]}} \frac{\partial h^{[2]}}{\partial g^{[2]}} \frac{\partial g^{[2]}}{\partial h^{[1]}} \frac{\partial h^{[1]}}{\partial g^{[1]}} \frac{\partial g^{[1]}}{\partial W^{[1]}}$$

$$(\hat{Y} - Y) \quad W^{[3]} \quad 0, 1 \quad W^{[2]} \quad 0, 1$$

1

$$\frac{\partial J}{\partial b^{[3]}} = \frac{\partial J}{\partial \hat{Y}} \frac{\partial \hat{Y}}{\partial g^{[3]}} \frac{\partial g^{[3]}}{\partial b^{[3]}}$$

1

$$\frac{\partial J}{\partial b^{[2]}} = \frac{\partial J}{\partial \hat{Y}} \frac{\partial \hat{Y}}{\partial g^{[3]}} \frac{\partial g^{[3]}}{\partial h^{[2]}} \frac{\partial h^{[2]}}{\partial g^{[2]}} \frac{\partial g^{[2]}}{\partial b^{[2]}}$$

1

$$\frac{\partial J}{\partial b^{[1]}} = \frac{\partial J}{\partial \hat{Y}} \frac{\partial \hat{Y}}{\partial g^{[3]}} \frac{\partial g^{[3]}}{\partial h^{[2]}} \frac{\partial h^{[2]}}{\partial g^{[2]}} \frac{\partial g^{[2]}}{\partial h^{[1]}} \frac{\partial h^{[1]}}{\partial g^{[1]}} \frac{\partial g^{[1]}}{\partial b^{[1]}}$$

$$(\hat{Y} - Y) \quad W^{[3]} \quad 0, 1 \quad W^{[2]} \quad 0, 1$$

Learning Flow

- Given a training dataset (\mathbf{X}, \mathbf{Y})
1. Initialize \mathbf{W}, \mathbf{b}
 2. Forward: calculate $\mathbf{h}^{[1]}, \mathbf{h}^{[2]}, \dots, \hat{\mathbf{Y}}, J$
 3. Backward: calculate $\frac{\partial J}{\partial \mathbf{W}}, \frac{\partial J}{\partial \mathbf{b}}$
 4. Optimize: e.g. SGD
 5. Verify with test dataset

in python keras

```
x_train,y_train,x_test,y_test=process_data()

model=Sequential()
model.add(Dense(10,input_dim=784,activation='relu'))
model.add(Dense(10,activation='softmax'))
model.compile(optimizer=SGD(lr=0.1),
              loss='categorical_crossentropy',
              metrics=['categorical_accuracy'])

traj=model.fit(x_train,
                y_train,
                epochs=50,
                batch_size=32,
                validation_data=(x_test, y_test),
                shuffle=True,verbose=0)

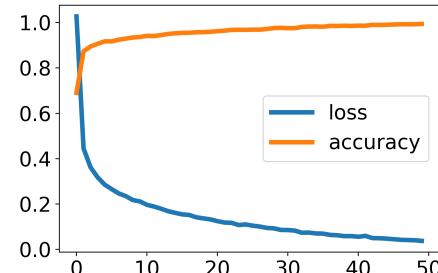
plt.plot(traj.history['loss'],linewidth=4,label='loss')
plt.plot(traj.history['categorical_accuracy'],linewidth=4,label='accuracy')
plt.legend()
```

x_train: (5000,784)
y_train: (5000,10) - one hot
x_test: (1000,784)
y_test: (1000,10) – one hot

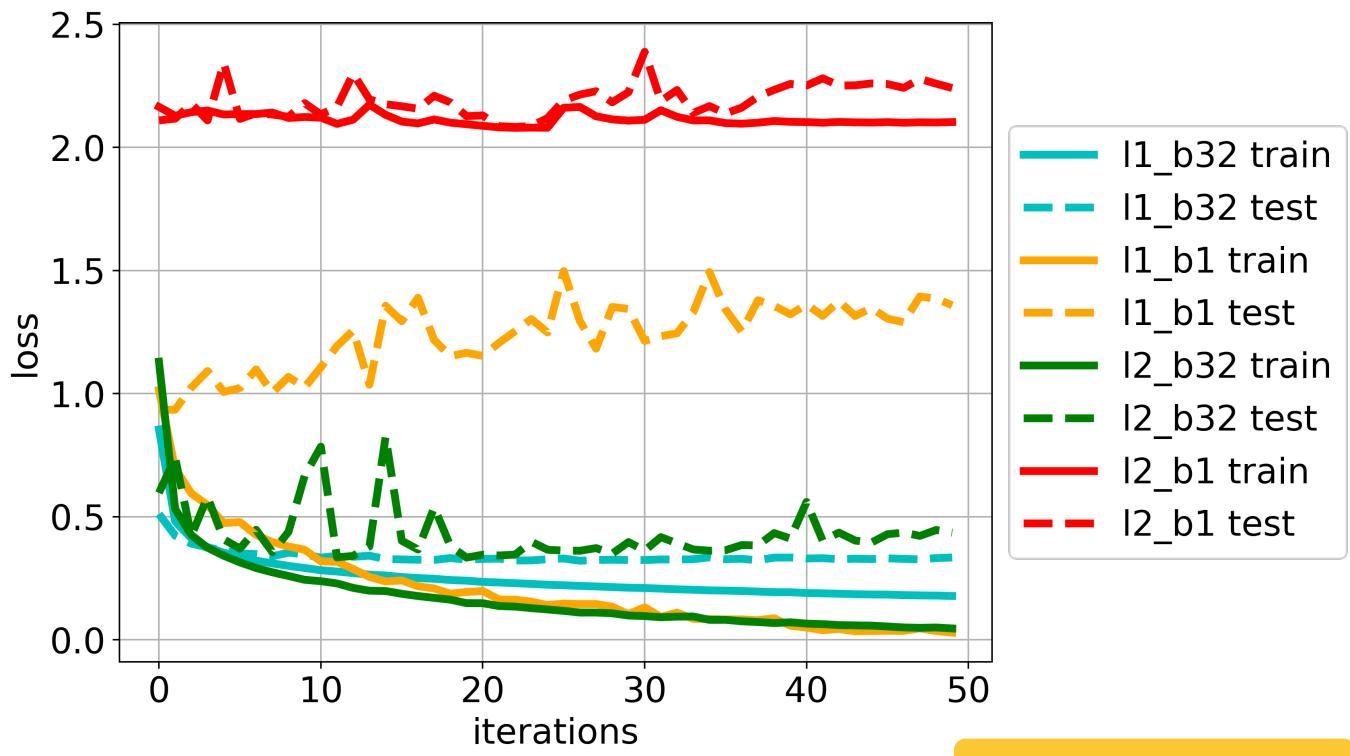
} Build model

} Training

Plot results



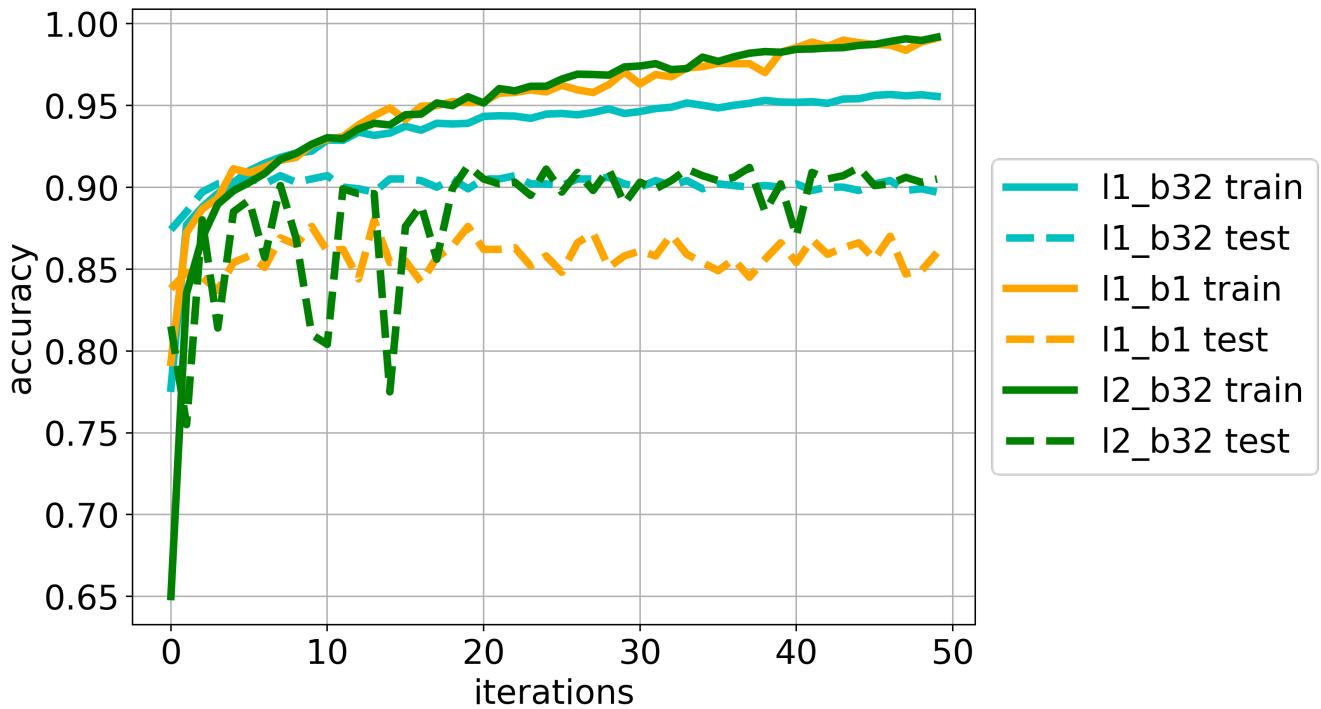
Learning behaviors [layers,batchsize]



Param size 7850

- l1_b32: one layer, batch size 32 <=> Softmax regression
- l1_b1: one layer, batch size 1 <=> Softmax regression SGD
- L2_b32: two layers, batch size 32 Param size 7960
- L2_b1: two layers, batch size 1 SGD

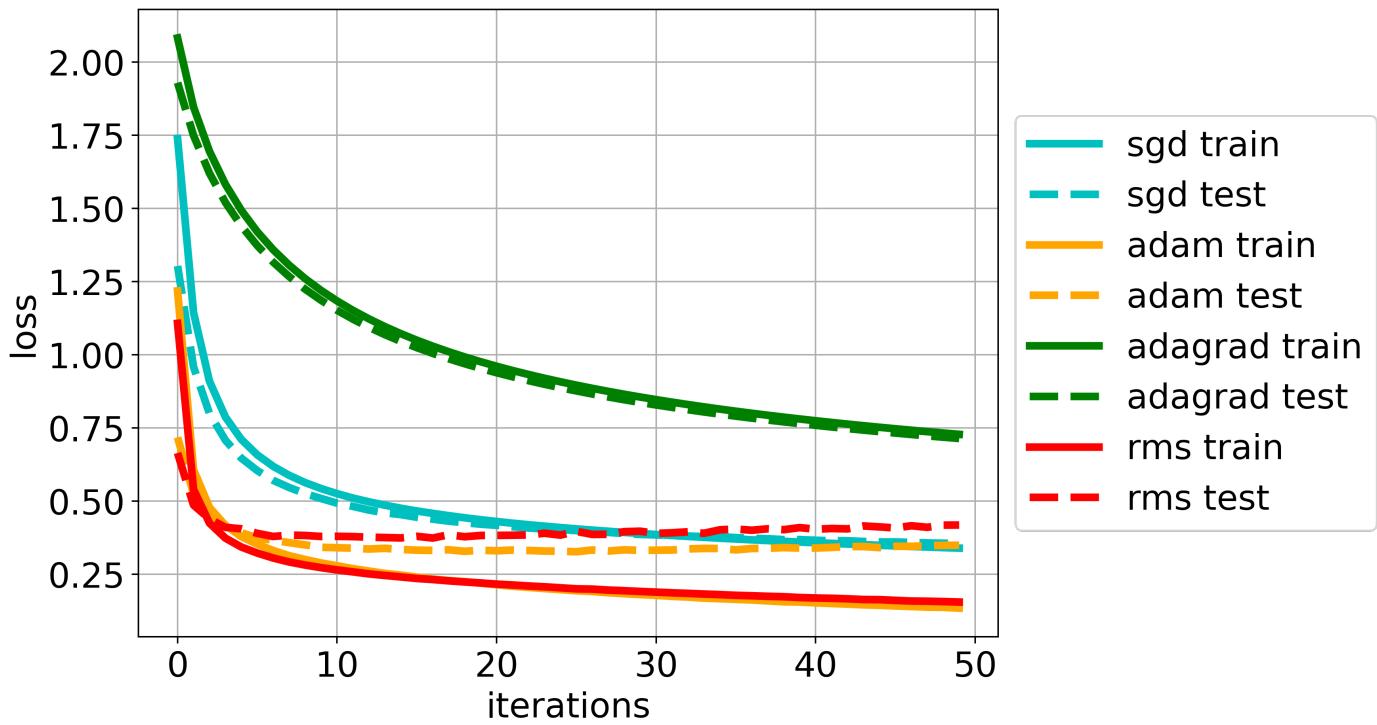
Learning behaviors [layers,batchsize]



- l_1_{b1} and l_2_{b32} fit the training set better than l_1_{b32}
- However, l_1_{b32} achieves stable generalization for the test set
- l_1_{b1} has the worst accuracy in test set

Behaviors of different optimizers

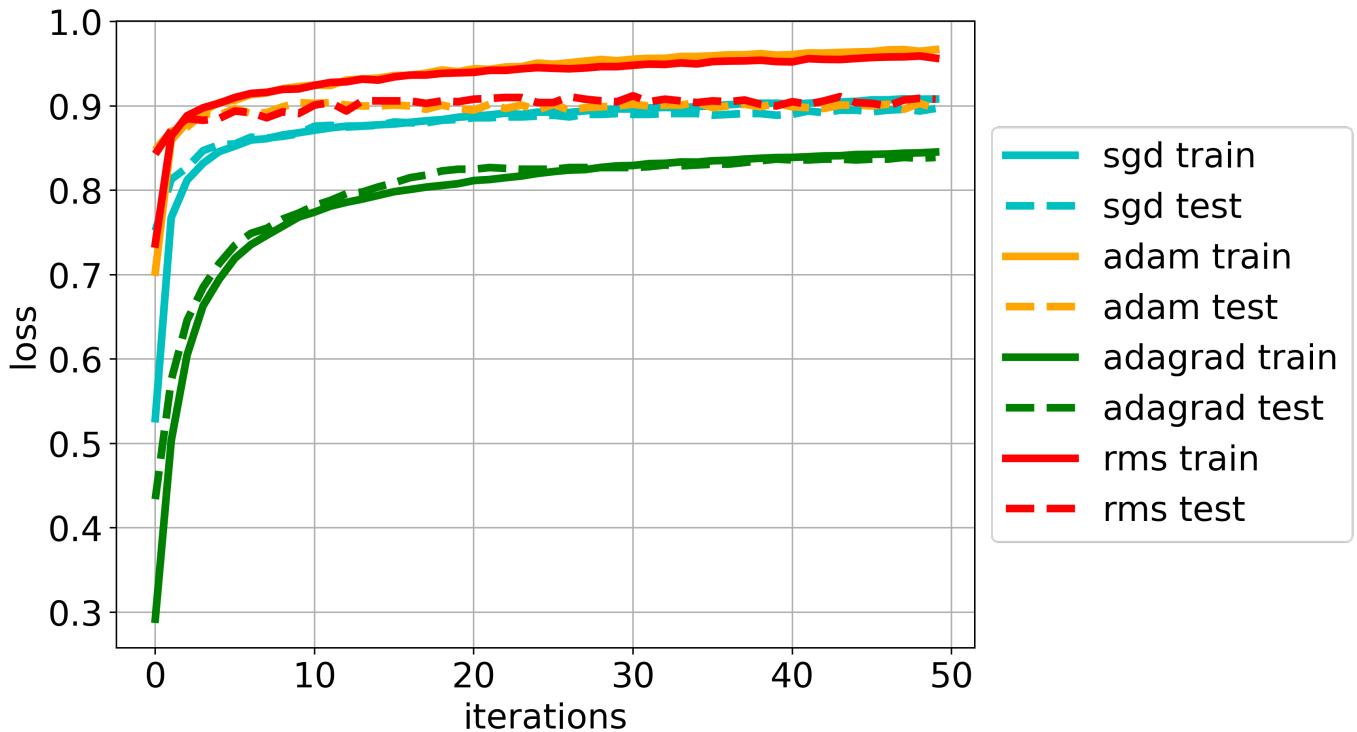
[layer : 1, batch: 32]



- Adam and RMSprop are relatively better than SGD and Adagrad

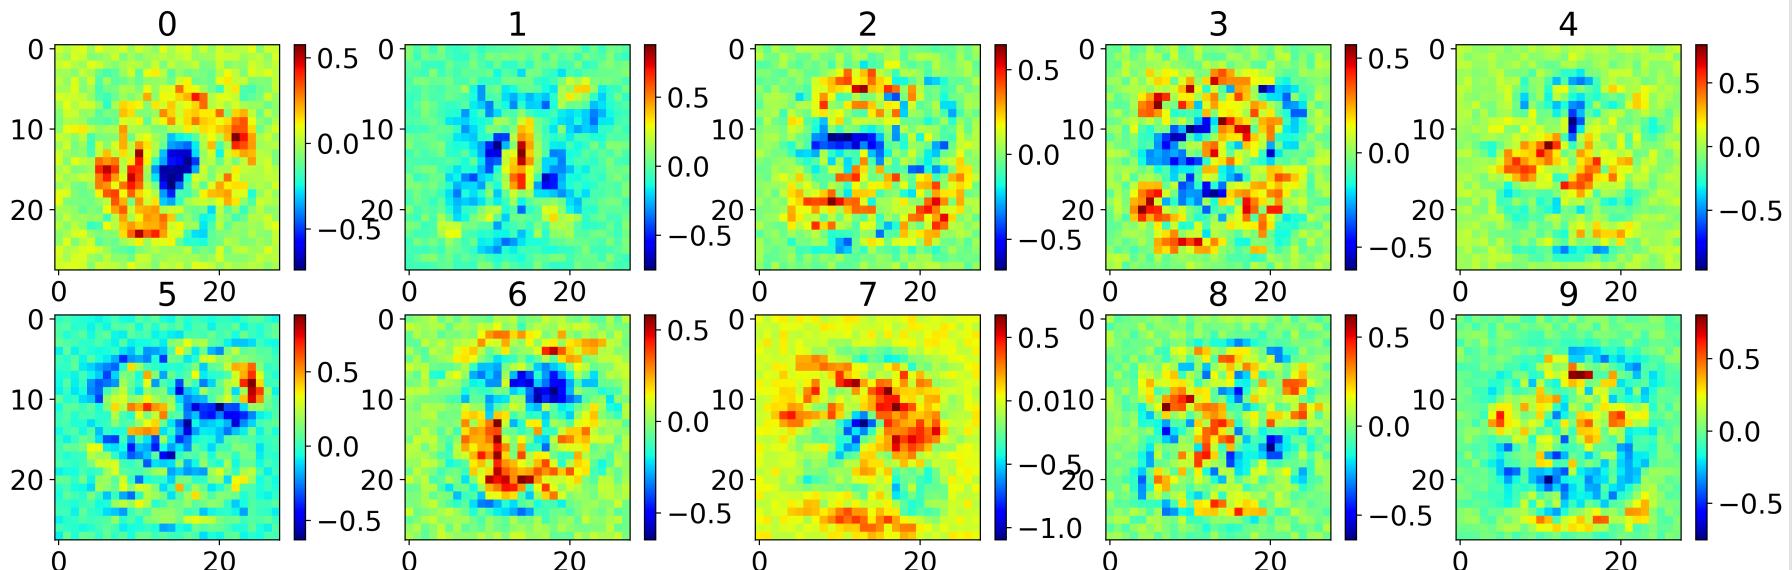
Behaviors of different optimizers

[layer : 1, batch: 32]



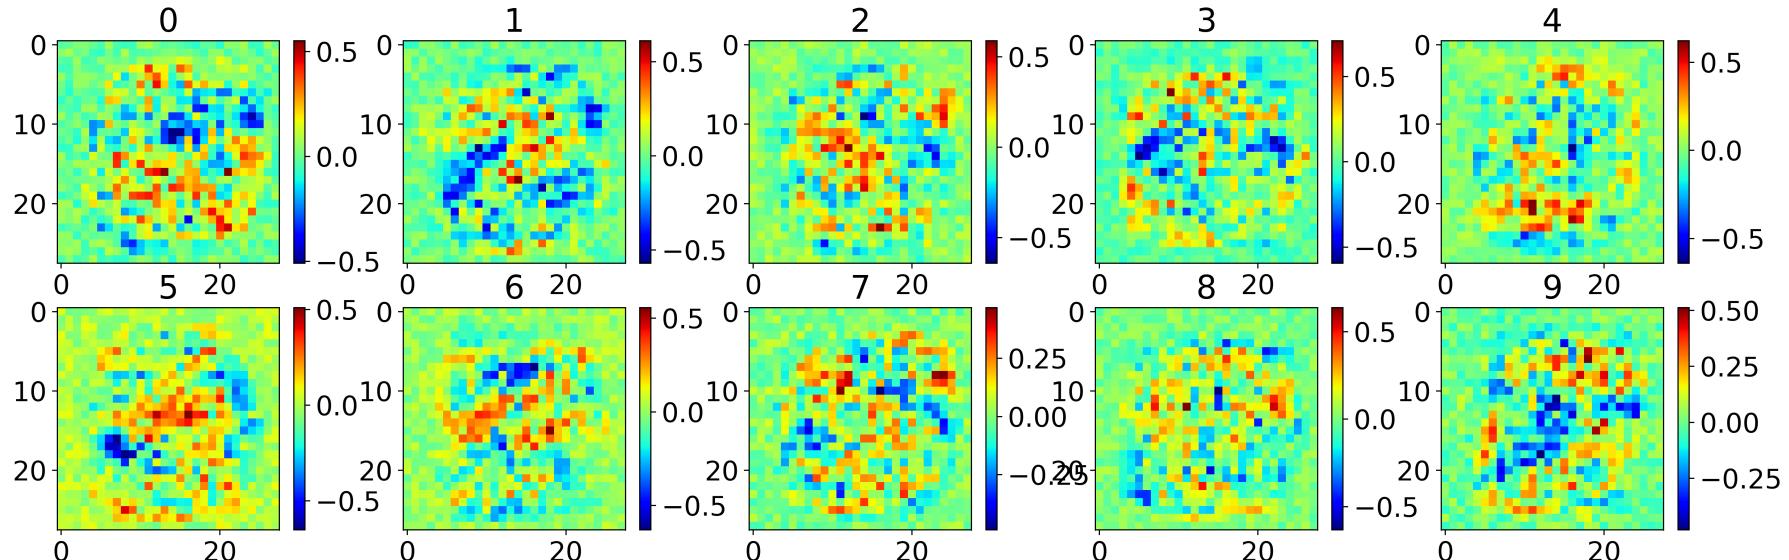
- Final accuracies on test sets are similar among SGD, Adam and RMSprop, however, Adam and RMSprop converge faster

Weights visualization for softmax regression (one-layer MLP)



Weights visualization for two-layer MLP

Layer 1



Layer 2

